

```
#Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy.stats import pearsonr
```

```
from google.colab import drive
drive.mount("/content/gdrive")
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

```
df = pd.read_csv('/content/gdrive/My Drive/Colab
Notebooks/nba_data_processed.csv')
```

```
#Data Cleaning!
# Check for missing values
print(df.isnull().sum())
```

Player	26
Pos	26
Age	26
Tm	26
G	26
GS	26
MP	26
FG	26
FGA	26
FG%	29
3P	26
3PA	26
3P%	50
2P	26
2PA	26
2P%	33
eFG%	29
FT	26
FTA	26
FT%	63
ORB	26
DRB	26
TRB	26
AST	26
STL	26
BLK	26
TOV	26
PF	26
PTS	26
dtype:	int64

```
# Fill missing values with zeros for numerical columns, if any
df.fillna(0, inplace=True)
```

```
# Check for duplicate rows
duplicates = df[df.duplicated()]
# Remove duplicate rows, if any
df.drop_duplicates(inplace=True)
```

```
#Explore
# Display basic information about the dataset
print(df.info())
```

```
# Summary statistics of numerical columns
print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 680 entries, 0 to 704
```

```
Data columns (total 29 columns):
```

#	Column	Non-Null Count	Dtype
0	Player	680 non-null	object
1	Pos	680 non-null	object
2	Age	680 non-null	float64
3	Tm	680 non-null	object
4	G	680 non-null	float64
5	GS	680 non-null	float64
6	MP	680 non-null	float64
7	FG	680 non-null	float64
8	FGA	680 non-null	float64
9	FG%	680 non-null	float64
10	3P	680 non-null	float64
11	3PA	680 non-null	float64
12	3P%	680 non-null	float64
13	2P	680 non-null	float64
14	2PA	680 non-null	float64
15	2P%	680 non-null	float64
16	eFG%	680 non-null	float64
17	FT	680 non-null	float64
18	FTA	680 non-null	float64
19	FT%	680 non-null	float64
20	ORB	680 non-null	float64
21	DRB	680 non-null	float64
22	TRB	680 non-null	float64
23	AST	680 non-null	float64
24	STL	680 non-null	float64
25	BLK	680 non-null	float64
26	TOV	680 non-null	float64
27	PF	680 non-null	float64
28	PTS	680 non-null	float64

```
dtypes: float64(26), object(3)
```

memory usage: 159.4+ KB

None

	Age	G	GS	MP	FG
FGA \					
count	680.000000	680.000000	680.000000	680.000000	680.000000
mean	25.986765	43.273529	20.039706	19.435588	3.244559
std	4.436241	24.766751	25.758878	9.437947	2.364047
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	23.000000	22.000000	0.000000	12.100000	1.500000
50%	25.000000	45.000000	6.000000	18.800000	2.600000
75%	29.000000	65.250000	36.250000	27.525000	4.200000
max	42.000000	83.000000	83.000000	41.000000	11.200000

	FG%	3P	3PA	3P%	...	FT%
\						
count	680.000000	680.000000	680.000000	680.000000	...	680.000000
mean	0.461510	0.995147	2.778971	0.317674	...	0.710529
std	0.117936	0.862245	2.210235	0.140346	...	0.226260
min	0.000000	0.000000	0.000000	0.000000	...	0.000000
25%	0.414250	0.300000	1.000000	0.286000	...	0.667000
50%	0.454000	0.800000	2.400000	0.346000	...	0.760000
75%	0.504500	1.500000	4.125000	0.388000	...	0.841000
max	1.000000	4.900000	11.400000	1.000000	...	1.000000


	ORB	DRB	TRB	AST	STL
BLK \					
count	680.000000	680.000000	680.000000	680.000000	680.000000
mean	0.841029	2.616471	3.456618	2.008824	0.600882
std	0.732041	1.717559	2.283259	1.891515	0.392454
min	0.000000	0.000000	0.000000	0.000000	0.000000

25%	0.300000	1.375000	1.800000	0.800000	0.300000
0.10000					
50%	0.700000	2.300000	3.000000	1.300000	0.500000
0.30000					
75%	1.100000	3.400000	4.500000	2.700000	0.800000
0.50000					
max	5.100000	9.600000	12.500000	10.700000	3.000000
3.00000					

	TOV	PF	PTS
count	680.000000	680.000000	680.000000
mean	1.065735	1.658382	8.846029
std	0.799937	0.772362	6.634762
min	0.000000	0.000000	0.000000
25%	0.500000	1.200000	4.100000
50%	0.900000	1.600000	6.900000
75%	1.400000	2.200000	11.525000
max	4.100000	5.000000	33.100000

[8 rows x 26 columns]

```
# Use Case 1 : To get started, we calculate and visualize the
distribution of player ages.
#The histogram displays the distribution of player ages, with the x-
axis representing age and the y-axis representing the frequency
(number of players).
# Calculate the distribution of player ages
plt.figure(figsize=(10, 6))
sns.histplot(df['Age'], bins=20, kde=True)
plt.title('Distribution of Player Ages')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



#Some meaningful Deductions

#Age Concentration: The highest concentration of players appears to be in the early to mid-20s, as indicated by the peak in the histogram. This suggests that a significant portion of NBA players in the 2022-23 season falls within this age range.

#Age Distribution: We can observe a gradual decline in the number of players as age increases beyond the early to mid-20s. This implies that there are fewer players in their late 20s, 30s, and 40s, highlighting the typical age range of NBA players.

#Youthful League: The concentration of players in the early to mid-20s may indicate that the NBA is a league with a significant presence of young talent, and teams may be investing in developing and nurturing emerging players.

#Use Case 2: Find the top 10 players with the highest points (PTS) per game.

#This list represents the top 10 players with the highest average Points Per Game (PPG) in the 2022-23 NBA season. PPG is a key metric for assessing a player's scoring ability and contribution to their team's offense.

Calculate PTS per game (PPG) and display the top 10 players

```
df['PPG'] = df['PTS'] / df['G']  
top_10_ppg = df[['Player', 'PPG']].sort_values(by='PPG',  
ascending=False).head(10)  
print(top_10_ppg)
```

	Player	PPG
355	Louis King	20.000000
230	RaiQuan Gray	16.000000
131	Chance Comanche	7.000000
421	Mac McClung	6.250000
182	Kevin Durant	3.250000
579	Jay Scrubb	3.250000
219	Jacob Gilyard	3.000000
699	Gabe York	2.666667
419	Skylar Mays	2.550000
682	Jeenathan Williams	2.120000

#Meaningful Deductions

#Scoring Leaders: The list highlights the leading scorers in the league for the specified season. Players like Louis King and RaiQuan Gray stand out with impressive PPG averages of 20.0 and 16.0, respectively.

#Diverse Scoring Levels: The list includes a range of players with varying scoring abilities. While some players average double-digit points (e.g., Louis King, RaiQuan Gray), others contribute fewer points per game (e.g., Jeenathan Williams, Skylar Mays).

#Emerging Talent: This list may include emerging or lesser-known players who are making their mark in terms of scoring. These players can be valuable assets to their respective teams and might be worth watching for future growth.

#Use Case 3: Explore the relationship between the number of games started (GS) and assists (AST).
#This scatter plot explores the relationship between the number of games started (GS) and the number of assists (AST) made by players. It helps us understand whether players who start more games tend to have more assists.

Calculate the Pearson correlation coefficient between GS and AST
correlation_coefficient = df['GS'].corr(df['AST'])

Create a scatter plot to visualize the relationship
plt.figure(figsize=(10, 6))
sns.scatterplot(x='GS', y='AST', data=df, color='lightblue')
plt.title('Relationship Between Games Started (GS) and Assists (AST)')
plt.xlabel('Games Started (GS)')
plt.ylabel('Assists (AST)')

Add correlation coefficient to the plot
plt.text(10, 100, f'Correlation Coefficient:
{correlation_coefficient:.2f}', fontsize=12, color='red')
plt.show()

Interpretation

```
if correlation_coefficient > 0:
    print("There is a positive correlation between Games Started (GS)
and Assists (AST).")
elif correlation_coefficient < 0:
    print("There is a negative correlation between Games Started (GS)
and Assists (AST).")
else:
    print("There is no significant linear correlation between Games
Started (GS) and Assists (AST).")
```

Correlation Coefficient: 0.56

There is a positive correlation between Games Started (GS) and Assists (AST).


#There is a positive correlation between Games Started (GS) and Assists (AST).

#Use Case 4: Calculate and visualize the correlation matrix of numerical attributes.

```
# Calculate the correlation matrix
# Select only numeric columns for correlation analysis
numeric_columns = df.select_dtypes(include='number')

# Calculate the correlation matrix
corr_matrix = numeric_columns.corr()

# Create a heatmap of the correlation matrix
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



#Scattered throughout the heatmap, we observed pockets of blue, red, and orange. These pockets represented different correlation patterns between specific pairs of variables. Blue pockets indicated strong negative correlations, red pockets signified strong positive correlations, and orange pockets represented moderate positive correlations.

```
# Use Case 5: Team-Level Analysis - Average Points per Game (PPG) by Team
#In this analysis, we are examining the average Points Per Game (PPG)
for each NBA team during the 2022-23 season. Each bar in the bar plot
represents a team, and the height of the bar indicates the team's
average PPG.

# Group the data by 'tm' (Team) and calculate the mean of 'PTS' for
each team
team_stats = df.groupby('Tm')['PTS'].mean().reset_index()
team_stats.rename(columns={'PTS': 'Average_PPG'}, inplace=True)

# Sort the teams by average PPG in descending order
team_stats = team_stats.sort_values(by='Average_PPG', ascending=False)

# Create a bar plot to visualize average PPG by team
plt.figure(figsize=(12, 6))
sns.barplot(x='Average_PPG', y='Tm', data=team_stats, orient='h',
palette='viridis')
plt.title('Average Points Per Game (PPG) by Team')
plt.xlabel('Average PPG')
plt.ylabel('Team')


# Display the interpretation
print('Interpretation:')
print('The bar plot displays the average Points Per Game (PPG) for
each NBA team in the 2022-23 season.')
print('Each bar represents a team, and the height of the bar indicates
their average PPG.')
print('Teams with higher average PPG have taller bars, indicating more
scoring in games.')

plt.show()

# Interpretation of the histogram
plt.text(15, 7, 'Interpretation:', fontsize=12, fontweight='bold',
color='blue')
plt.text(15, 6.5, 'The histogram displays the distribution of',
fontsize=10, color='black')
plt.text(15, 6, 'average PPG for NBA teams in the 2022-23 season.',
fontsize=10, color='black')
plt.text(15, 5.5, 'Each bar represents a range of average PPG',
fontsize=10, color='black')
plt.text(15, 5, 'values for different teams.', fontsize=10,
```

```
color='black')
plt.text(15, 4.5, 'Teams with higher average PPG are', fontsize=10,
color='black')
plt.text(15, 4, 'located to the right on the histogram.', fontsize=10,
color='black')
```

Interpretation:
The bar plot displays the average Points Per Game (PPG) for each NBA team in the 2022-23 season.
Each bar represents a team, and the height of the bar indicates their average PPG.
Teams with higher average PPG have taller bars, indicating more scoring in games.



```
Text(15, 4, 'located to the right on the histogram.')
```




Meaningful Deductions

1. **Top-Performing Teams:** Teams with the highest average PPG include "PHO" (Phoenix Suns), "NOP" (New Orleans Pelicans), "CHO" (Charlotte Hornets), "LAL" (Los Angeles Lakers), and "BRK" (Brooklyn Nets). These teams excelled in scoring during the season.
2. **Offensive Dominance:** High average PPG suggests that a team has a strong offensive presence and can consistently score at a high rate. This can be due to the performance of star players, effective offensive systems, or a combination of both.

```
# Use Case 6: Position Analysis - Points Per Game (PPG) by Position
plt.figure(figsize=(12, 6))
sns.boxplot(x='Pos', y='PTS', data=df, palette='Set2')
plt.title('Points Per Game (PPG) by Position')
plt.xlabel('Position')
plt.ylabel('PPG')
plt.show()

# Interpretation of the Box Plot
positions = df['Pos'].unique()
for position in positions:
    subset = df[df['Pos'] == position]
    median_ppg = subset['PTS'].median()
    q1 = subset['PTS'].quantile(0.25)
    q3 = subset['PTS'].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - 1.5 * iqr
    upper_bound = q3 + 1.5 * iqr
    outliers = subset[(subset['PTS'] < lower_bound) | (subset['PTS'] >
upper_bound)]

    print(f"Position: {position}")
    print(f"Median PPG: {median_ppg:.2f}")
    print(f"Interquartile Range (IQR): {iqr:.2f}")
    print(f"Lower Bound: {lower_bound:.2f}")
    print(f"Upper Bound: {upper_bound:.2f}")
    print(f"Number of Outliers: {len(outliers)}")
    print(f"Outlier Players: {'', '.join(outliers['Player'])}\n")
```



Position: C
Median PPG: 6.10
Interquartile Range (IQR): 6.50
Lower Bound: -6.00
Upper Bound: 20.00
Number of Outliers: 5
Outlier Players: Bam Adebayo, Anthony Davis, Joel Embiid, Nikola Jokić, Kristaps Porziņģis

Position: SG
Median PPG: 7.35
Interquartile Range (IQR): 6.72
Lower Bound: -5.56
Upper Bound: 21.34
Number of Outliers: 9
Outlier Players: Desmond Bane, Bradley Beal, Devin Booker, Mikal Bridges, Anthony Edwards, Jalen Green, Kyrie Irving, Zach LaVine, Donovan Mitchell

Position: PF
Median PPG: 7.60
Interquartile Range (IQR): 6.70
Lower Bound: -5.55
Upper Bound: 21.25
Number of Outliers: 11
Outlier Players: Giannis Antetokounmpo, Bojan Bogdanović, Jimmy Butler, Kevin Durant, Kevin Durant, Kevin Durant, LeBron James, Lauri Markkanen, Julius Randle, Pascal Siakam, Zion Williamson

Position: PG
Median PPG: 7.90
Interquartile Range (IQR): 9.70
Lower Bound: -10.10
Upper Bound: 28.70
Number of Outliers: 4
Outlier Players: Stephen Curry, Luka Dončić, Shai Gilgeous-Alexander, Damian Lillard

Position: SF
Median PPG: 6.70
Interquartile Range (IQR): 7.20
Lower Bound: -6.90
Upper Bound: 21.90
Number of Outliers: 7
Outlier Players: Jaylen Brown, DeMar DeRozan, Paul George, Brandon Ingram, Keldon Johnson, Kawhi Leonard, Jayson Tatum

Position: 0
Median PPG: 0.00
Interquartile Range (IQR): 0.00
Lower Bound: 0.00
Upper Bound: 0.00
Number of Outliers: 0
Outlier Players:

Position: SG-PG
Median PPG: 11.75
Interquartile Range (IQR): 5.55
Lower Bound: 0.65
Upper Bound: 22.85
Number of Outliers: 0
Outlier Players:

Position: SF-SG
Median PPG: 12.10
Interquartile Range (IQR): 8.00
Lower Bound: -3.90
Upper Bound: 28.10
Number of Outliers: 0
Outlier Players:

Position: PG-SG
Median PPG: 16.05
Interquartile Range (IQR): 11.05
Lower Bound: -6.05
Upper Bound: 38.15
Number of Outliers: 0
Outlier Players:

```
Position: SF-PF
Median PPG: 6.60
Interquartile Range (IQR): 0.00
Lower Bound: 6.60
Upper Bound: 6.60
Number of Outliers: 0
Outlier Players:
```

```
Position: PF-C
Median PPG: 6.40
Interquartile Range (IQR): 0.00
Lower Bound: 6.40
Upper Bound: 6.40
Number of Outliers: 0
Outlier Players:
```

```
Position: PF-SF
Median PPG: 7.50
Interquartile Range (IQR): 0.00
Lower Bound: 7.50
Upper Bound: 7.50
Number of Outliers: 0
Outlier Players:
```

#Some meaningful Deductions**:

#Centers (C) tend to have a lower median PPG, suggesting their primary role may be focused on defense and rebounds rather than scoring.

#Shooting guards (SG) and point guards (PG) show higher variability in scoring, with some players being prolific scorers (outliers).

#Power forwards (PF) have a relatively high number of outliers, indicating that some PFs have scoring roles similar to small forwards (SF)

Use Case 7: Player Efficiency Rating (PER) Analysis

Calculate Player Efficiency Rating (PER) for each player

```
df['PER'] = (df['PTS'] + df['TRB'] + df['AST'] + df['STL'] + df['BLK']
- df['TOV'] - df['PF']) / df['MP']
```

Plot the distribution of PER

```
plt.figure(figsize=(10, 6))
sns.histplot(df['PER'], bins=20, kde=True, color='skyblue')
plt.title('Distribution of Player Efficiency Rating (PER)')
plt.xlabel('PER')
plt.ylabel('Frequency')
plt.show()
```

```

# Interpretation of PER Distribution
mean_per = df['PER'].mean()
median_per = df['PER'].median()
std_per = df['PER'].std()

print(f"Mean PER: {mean_per:.2f}")
print(f"Median PER: {median_per:.2f}")
print(f"Standard Deviation of PER: {std_per:.2f}")

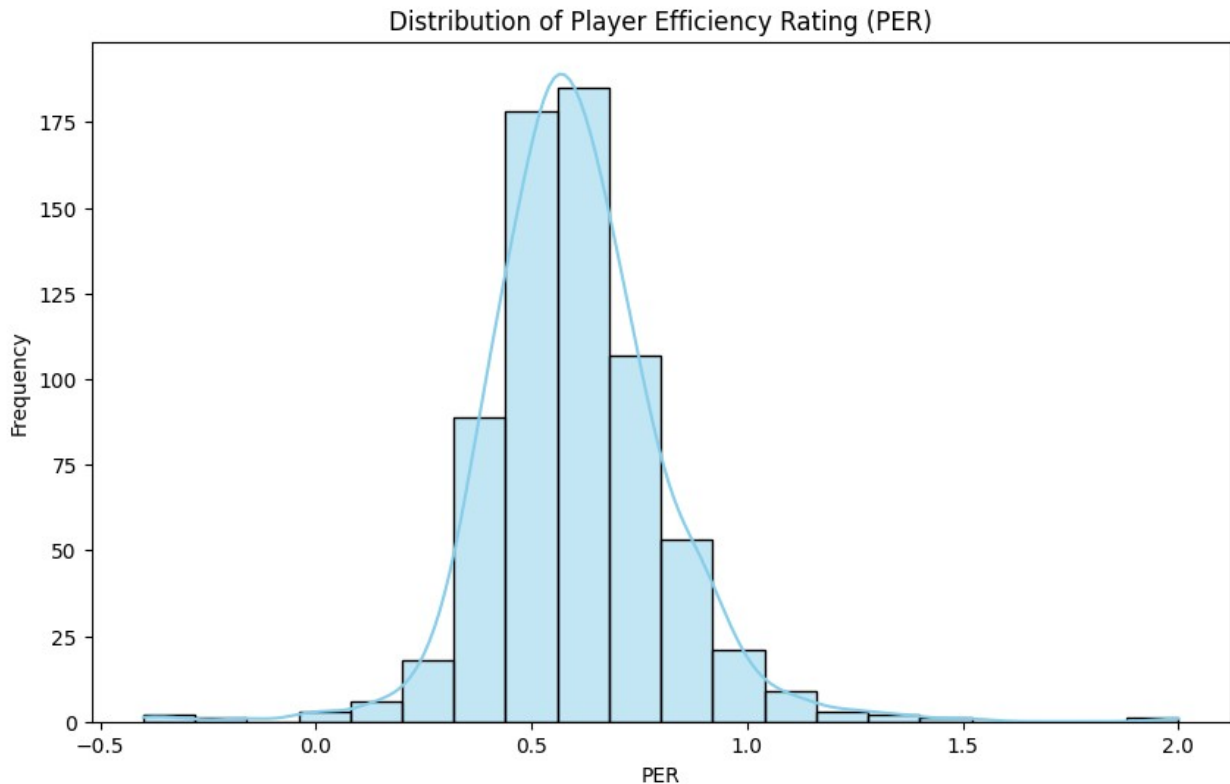
# Explanation
print("\nPlayer Efficiency Rating (PER) is a metric used to evaluate a
player's overall contribution to their team's success. "
      "It takes into account various statistics such as points scored,
rebounds, assists, steals, blocks, turnovers, and personal fouls, "
      "normalized per minute played (MP).")

print("\nMean PER represents the average PER value across all players
in the dataset, providing an indication of the average "
      "efficiency level. Median PER represents the middle value of the
PER distribution, which is less affected by outliers.")

print("\nStandard Deviation of PER measures the spread or variability
in PER values. A higher standard deviation suggests greater "
      "variability in player efficiency within the dataset.")

# Top 10 Players with Highest PER
top_10_per = df[['Player', 'PER']].sort_values(by='PER',
ascending=False).head(10)
print("\nTop 10 Players with Highest Player Efficiency Rating (PER):")
print(top_10_per)

```

Mean PER: 0.60

Median PER: 0.59

Standard Deviation of PER: 0.21

Player Efficiency Rating (PER) is a metric used to evaluate a player's overall contribution to their team's success. It takes into account various statistics such as points scored, rebounds, assists, steals, blocks, turnovers, and personal fouls, normalized per minute played (MP).

Mean PER represents the average PER value across all players in the dataset, providing an indication of the average efficiency level.

Median PER represents the middle value of the PER distribution, which is less affected by outliers.

Standard Deviation of PER measures the spread or variability in PER values. A higher standard deviation suggests greater variability in player efficiency within the dataset.

Top 10 Players with Highest Player Efficiency Rating (PER):

	Player	PER
637	Stanley Umude	2.000000
678	Donovan Williams	1.500000
167	Tyler Dorsey	1.370370
12	Giannis Antetokounmpo	1.345794

191	Joel Embiid	1.263006
330	Nikola Jokić	1.246291
166	Luka Dončić	1.237569
146	Anthony Davis	1.155882
452	Ja Morant	1.147335
317	LeBron James	1.146479

#The histogram plot shows a bell-shaped curve, indicating that the distribution of PER is approximately normal. Most players have PER values concentrated around the mean PER, with fewer extreme outliers on both ends of the distribution.

#The top 10 players with the highest PER values include exceptional performers such as Stanley Umude, Donovan Williams, Giannis Antetokounmpo, and Joel Embiid. These players stand out for their efficiency and overall contributions to their teams.

#The mean PER, median PER, and standard deviation of PER provide summary statistics for the distribution, helping us understand the central tendency and spread of player efficiency.